

MUSTHOFA JOKO ANGGORO

✉ musthofaja.topa@gmail.com | [linkedin.com/in/musthofa-joko-anggoro](https://www.linkedin.com/in/musthofa-joko-anggoro) | github.com/topahilangharapan | portfolio-topahilangharapan-personal.vercel.app

April 29, 2026

Hiring Team
StraitsX
Singapore

Dear StraitsX Hiring Team,

Reducing manual booking intervention by 85% wasn't the original goal — building an automation pipeline that operators could actually trust was. That same instinct is what draws me to StraitsX's AI Taskforce: the role isn't about using AI tools, it's about designing the agentic workflows and internal infrastructure that make those tools genuinely useful at scale.

At Traveloka, one of the largest flight inventory providers had no automated reschedule path — operators manually processed thousands of booking modifications daily. I built an end-to-end pipeline with configurable scrapers, failover backup mechanisms, and 99.9% processing accuracy, cutting manual intervention by 85%. The team stopped doing repetition and started handling exceptions. Alongside that, I integrated GitHub Copilot CLI into our development workflow to compress feedback loops and raise automated test coverage — the AI didn't write the features, but it meaningfully changed how fast we could ship them safely.

The part of this role I find most compelling is the MCP layer — the connective tissue between AI agents and internal systems. That's where the real adoption problems live: not whether the model is capable, but whether it has the right context at the right time. I've been exploring this both at work and outside it — which is also how I ended up building agent configuration and authoring workflows into my own portfolio. The governance piece matters too: a system that can't be audited or safely constrained isn't a productivity tool, it's a liability. That's a lesson I already applied when I built a Credential Management System with AES-GCM encryption and granular audit logging at Traveloka.

I'd welcome the chance to talk through how StraitsX's Taskforce is thinking about the boundary between agent capability and safety governance — that's the design question I find most interesting right now, and I suspect it's the hardest one on the team's roadmap.

Sincerely,

Musthofa Joko Anggoro

University of Indonesia, Fasilkom · Information Systems · GPA 3.84/4.00 · Expected graduation June–August 2026