

MUSTHOFA JOKO ANGGORO

✉ musthofaja.topa@gmail.com |  linkedin.com/in/musthofa-joko-anggoro | 
github.com/topahilangharapan |  portfolio-topahilangharapan-personal.vercel.app

April 30, 2026

Crossing Hurdles Hiring Team
Crossing Hurdles

Re: Software Engineer — Remote (Contract)

Building reliable automation that processes thousands of transactions daily with 99.9% accuracy is one thing — integrating that same reliability mindset into AI agent pipelines is the problem space I've been moving toward. At Traveloka, I didn't just use LLM-assisted tooling (GitHub Copilot CLI) to ship code faster; I started thinking about what it takes to make AI workflows production-grade: deterministic failover, observable state transitions, and outcomes you can actually measure. When Crossing Hurdles describes designing scalable features across Python, React, GoLang, and Node.js for high-performance AI agent applications, that's the exact intersection of backend engineering and LLM integration I want to go deeper in.

The clearest example of building at that intersection was a reschedule automation pipeline I engineered at Traveloka. One of the company's largest flight inventory providers had no automated modification path — operators manually handled over a thousand booking changes daily, with all the error surface that implies. I designed an end-to-end pipeline with configurable scrapers, intelligent failover logic, and idempotent state transitions, cutting manual intervention by **85%** while sustaining **99.9% processing accuracy**. The hard part wasn't the automation itself — it was making the system observable and trustworthy enough that engineers downstream could rely on it without second-guessing every output. That's the same bar I'd hold AI agent workflows to.

Crossing Hurdles' focus on advanced AI agent training and LLM workflow optimization is where I find the most intellectually interesting engineering right now. I've spent the last year actively exploring the gap between model capability and production reliability — messing with how LLMs behave in agentic loops, where they fail silently, and how you instrument them to catch those failures early. The polyglot stack (Python, GoLang, Node.js, React) maps directly to what I work in, and the cloud infrastructure side — deploying and maintaining systems on AWS — is the direction I've been building toward from a GCP and Kubernetes foundation at Traveloka.

I'd welcome the chance to talk through how Crossing Hurdles thinks about reliability and observability for AI agent workflows — and whether my background in backend automation, LLM tooling, and distributed systems maps to where the team needs to grow. I'm based in Indonesia, fully remote, and available to start on a contract basis.

Musthofa Joko Anggoro